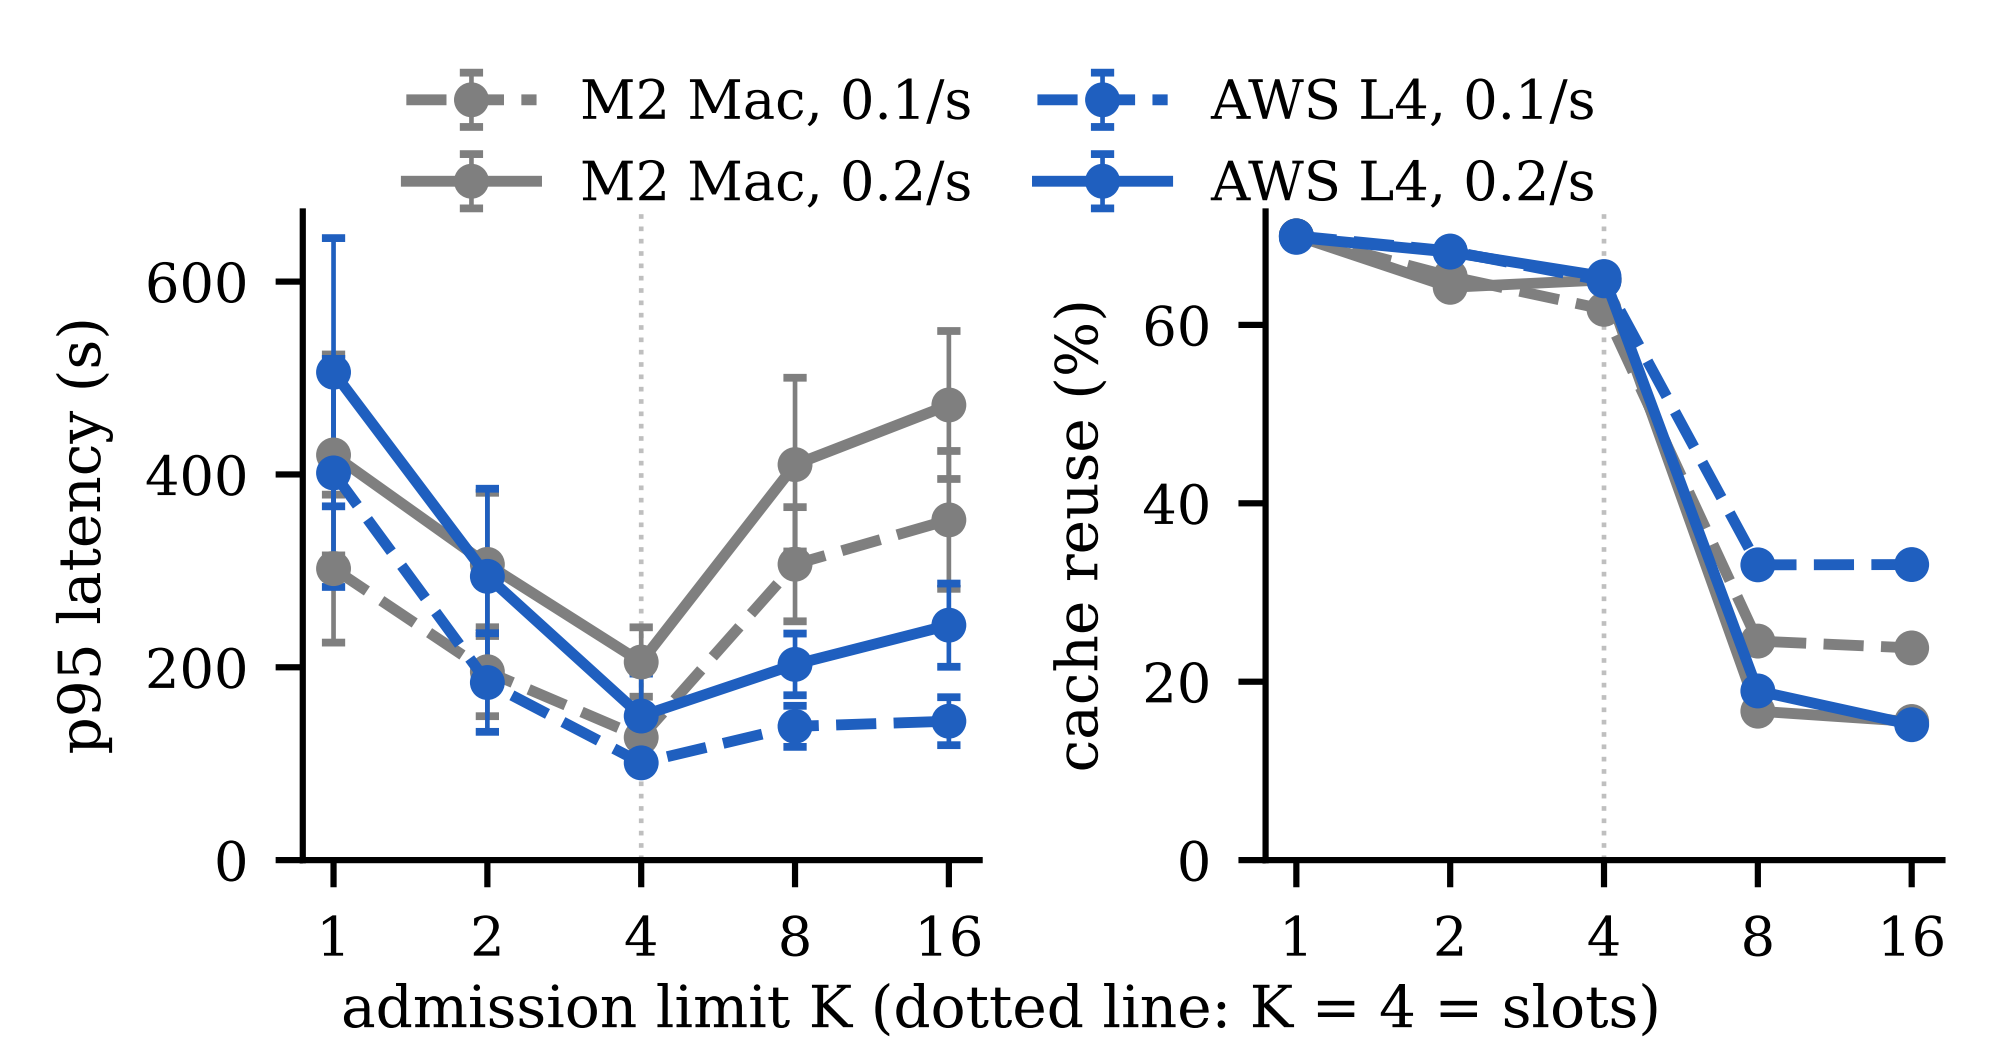
M2 Mac, 0.1/s
M2 Mac, 0.2/s
AWS L4, 0.1/s
AWS L4, 0.2/s
p95 latency (s)
0
200
400
600
cache reuse (%)
0
20
40
60
1
2
4
8
16
admission limit K (dotted line: K = 4 = slots)